

Mushroom Classification

Unsupervised Machine Learning

Table of contents

1	Introduction	2
1.1	Exploratory Analysis	7
1.2	Data Preprocessing for Unsupervised Learning .	13
2	Unsupervised Learning Analysis	14
2.1	Principal Component Analysis (PCA)	14
2.2	K-means Clustering	17
2.3	Matrix Factorization (Non-negative Matrix Factorization)	20
2.4	Hierarchical Clustering	28
2.5	t-SNE Visualization	33
2.6	Feature Importance through Clustering	34
2.7	Association Analysis	35
2.8	Cluster Validation	38
3	Conclusion	39

```
suppressPackageStartupMessages(library(tidyverse))
suppressPackageStartupMessages(library(ggplot2))
suppressPackageStartupMessages(library(reshape2))
suppressPackageStartupMessages(library(cluster))
suppressPackageStartupMessages(library(factoextra))
suppressPackageStartupMessages(library(NbClust))
suppressPackageStartupMessages(library(corrplot))
```

```
suppressPackageStartupMessages(library(FactoMineR))
suppressPackageStartupMessages(library(VIM))
suppressPackageStartupMessages(library(NMF))
suppressPackageStartupMessages(library(Rtsne))
suppressPackageStartupMessages(library(umap))
suppressPackageStartupMessages(library(pheatmap))
set.seed(123)
```

1 Introduction

In this project, we'll explore the Mushroom dataset using unsupervised machine learning techniques to discover hidden patterns and natural groupings within mushroom characteristics. Rather than predicting edibility directly, we'll investigate whether mushrooms naturally cluster into distinct groups based on their physical features, and examine how these discovered patterns relate to the known edible/poisonous classifications.

The dataset includes artificial descriptions of various mushrooms with their characteristics, based upon the Audobon Society Field Guide. Using unsupervised learning, we'll attempt to uncover the underlying structure in mushroom features without using the edibility labels during analysis. For more information about this data, it can be found at the [UC Irvine Machine Learning Repository](#).

First, we need to load in our data.

```
data <- read.csv("data/mushrooms.csv")

head(data)
```

	class	cap.shape	cap.surface	cap.color	bruises	odor	gill.attachment
1	p	x	s	n	t	p	f
2	e	x	s	y	t	a	f
3	e	b	s	w	t	l	f
4	p	x	y	w	t	p	f
5	e	x	s	g	f	n	f

	e	x	y	y	t	a	f
	gill.spacing gill.size gill.color stalk.shape stalk.root						
1		c	n	k	e	e	
2		c	b	k	e	c	
3		c	b	n	e	c	
4		c	n	n	e	e	
5		w	b	k	t	e	
6		c	b	n	e	c	
	stalk.surface.above.ring stalk.surface.below.ring stalk.color.above.ring						
1			s		s		w
2			s		s		w
3			s		s		w
4			s		s		w
5			s		s		w
6			s		s		w
	stalk.color.below.ring veil.type veil.color ring.number ring.type						
1			w	p	w	o	p
2			w	p	w	o	p
3			w	p	w	o	p
4			w	p	w	o	p
5			w	p	w	o	e
6			w	p	w	o	p
	spore.print.color population habitat						
1		k		s	u		
2		n		n	g		
3		n		n	m		
4		k		s	u		
5		n		a	g		
6		k		n	g		

```
summary(data)
```

class	cap.shape	cap.surface	cap.color
Length:8124	Length:8124	Length:8124	Length:8124
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
bruises	odor	gill.attachment	gill.spacing
Length:8124	Length:8124	Length:8124	Length:8124
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character

gill.size	gill.color	stalk.shape	stalk.root
Length:8124	Length:8124	Length:8124	Length:8124
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
stalk.surface.above.ring	stalk.surface.below.ring	stalk.color.above.ring	
Length:8124	Length:8124	Length:8124	
Class :character	Class :character	Class :character	
Mode :character	Mode :character	Mode :character	
stalk.color.below.ring	veil.type	veil.color	
Length:8124	Length:8124	Length:8124	
Class :character	Class :character	Class :character	
Mode :character	Mode :character	Mode :character	
ring.number	ring.type	spore.print.color	population
Length:8124	Length:8124	Length:8124	Length:8124
Class :character	Class :character	Class :character	Class :character
Mode :character	Mode :character	Mode :character	Mode :character
habitat			
Length:8124			
Class :character			
Mode :character			

Before we get too far, we should make sure we know what our column names are.

```
colnames(data)
```

```
[1] "class"
[3] "cap.surface"
[5] "bruises"
[7] "gill.attachment"
[9] "gill.size"
[11] "stalk.shape"
[13] "stalk.surface.above.ring"
[15] "stalk.color.above.ring"
[17] "veil.type"
[19] "ring.number"
[21] "spore.print.color"
[23] "habitat"

"cap.shape"
"cap.color"
"odor"
"gill.spacing"
"gill.color"
"stalk.root"
"stalk.surface.below.ring"
"stalk.color.below.ring"
"veil.color"
"ring.type"
"population"
```

Let's examine the distribution of our known classes (which we'll use later for validation but not during clustering):

```
distribution <- table(data$class)
distribution
```

```
      e      p
4208 3916
```

```
distribution[1] / (distribution[1] + distribution[2])
```

```
      e
0.5179714
```

```
distribution[2] / (distribution[1] + distribution[2])
```

```
      p
0.4820286
```

This shows a pretty even split. Let's also check for any missing values in the data.

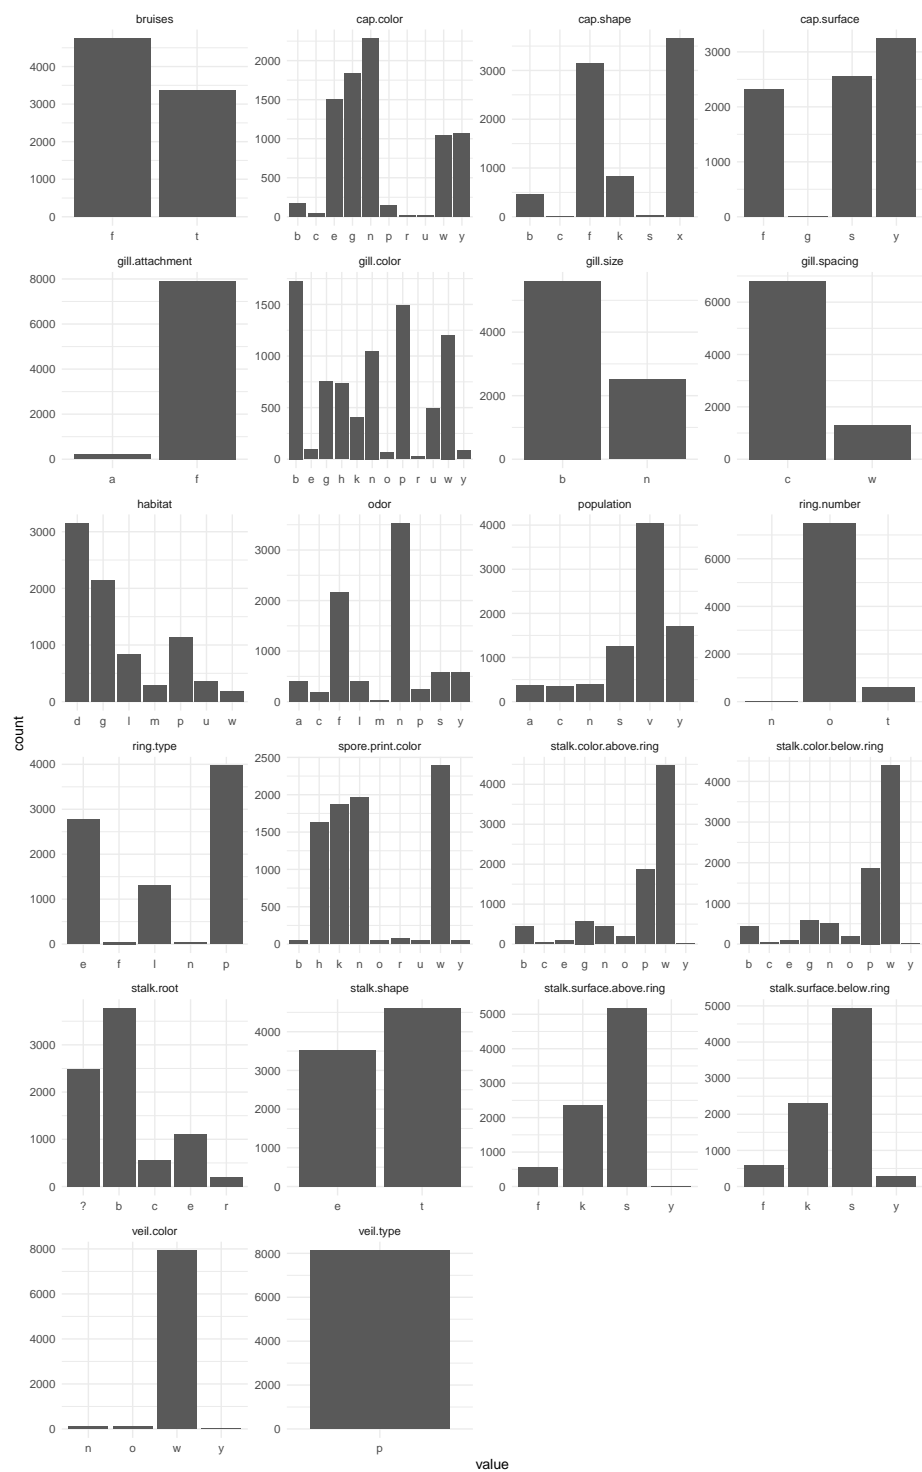
```
anyNA(data)
```

```
[1] FALSE
```

We can see that there are no missing values in this dataset. Now let's visualize the distribution of features:

```
data_long <- data %>% gather(data, "key", class:habitat)
colnames(data_long) <- c("group", "value")
data_long <- data_long[data_long$group != "class",]

ggplot(data_long) +
  geom_bar(aes(x = value, group = group)) +
  facet_wrap(~group, ncol = 4, scales = "free") +
  theme_minimal()
```

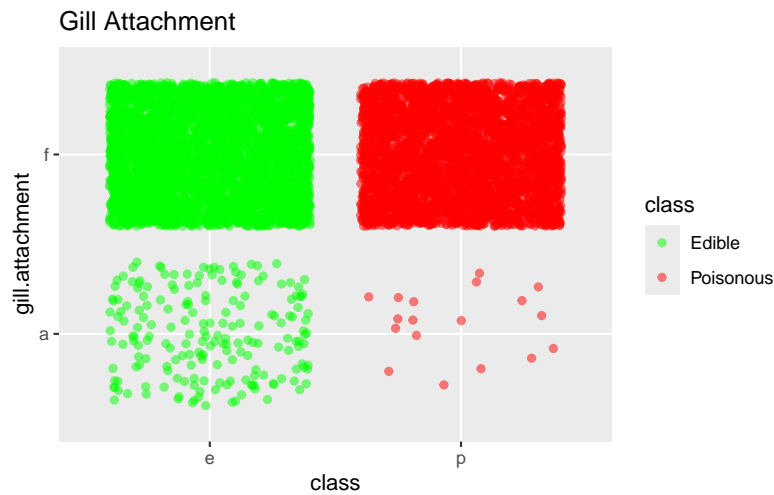


- cap-shape: bell=b, conical=c, convex=x, flat=f, knobbed=k, sunken=s
- cap-surface: fibrous=f, grooves=g, scaly=y, smooth=s
- cap-color: brown=n, buff=b, cinnamon=c, gray=g, green=r, pink=p, purple=u, red=e, white=w, yellow=y
- bruises: bruises=t, no=f
- odor: almond=a, anise=l, creosote=c, fishy=y, foul=f, musty=m, none=n, pungent=p, spicy=s
- gill-attachment: attached=a, descending=d, free=f, notched=n
- gill-spacing: close=c, crowded=w, distant=d
- gill-size: broad=b, narrow=n
- gill-color: black=k, brown=n, buff=b, chocolate=h, gray=g, green=r, orange=o, pink=p, purple=u, red=e, white=w, yellow=y
- stalk-shape: enlarging=e, tapering=t
- stalk-root: bulbous=b, club=c, cup=u, equal=e, rhizomorphs=z, rooted=r, missing=?
- stalk-surface-above-ring: fibrous=f, scaly=y, silky=k, smooth=s
- stalk-surface-below-ring: fibrous=f, scaly=y, silky=k, smooth=s
- stalk-color-above-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
- stalk-color-below-ring: brown=n, buff=b, cinnamon=c, gray=g, orange=o, pink=p, red=e, white=w, yellow=y
- veil-type: partial=p, universal=u
- veil-color: brown=n, orange=o, white=w, yellow=y
- ring-number: none=n, one=o, two=t
- ring-type: cobwebby=c, evanescent=e, flaring=f, large=l, none=n, pendant=p, sheathing=s, zone=z
- spore-print-color: black=k, brown=n, buff=b, chocolate=h, green=r, orange=o, purple=u, white=w, yellow=y
- population: abundant=a, clustered=c, numerous=n, scattered=s, several=v, solitary=y
- habitat: grasses=g, leaves=l, meadows=m, paths=p, urban=u, waste=w, woods=d

1.1 Exploratory Analysis

We can also see that some of our variables are predominantly one value. Let's take a closer look at these predictors to see if they might have an impact on our model if it is worth removing them in order to prevent an overly weighted predictor from impacting the model.

```
ggplot(data, aes(x = class, y = gill.attachment, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Gill Attachment") +
  scale_color_manual(breaks = c("e", "p"),
                    labels = c("Edible", "Poisonous"),
                    values = c("green", "red"))
```




```
table(data$class, data$gill.attachment)
```

	a	f
e	192	4016
p	18	3898

Here we can see a pretty even split between the edible and poisonous classes. There appears to be a tendency for attached veils to be more often edible rather than poisonous

```
ggplot(data, aes(x = class, y = gill.spacing, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Gill Attachment") +
  scale_color_manual(breaks = c("e", "p"),
                     labels = c("Edible", "Poisonous"),
                     values = c("green", "red"))
```





```
table(data$class, data$gill.spacing)
```

	c	w
e	3008	1200
p	3804	112

This is also pretty split when considering gill spacing. Here we can see a tendency for crowded gills to be edible more often.

```
ggplot(data, aes(x = class, y = ring.number, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Gill Attachment") +
  scale_color_manual(breaks = c("e", "p"),
                    labels = c("Edible", "Poisonous"),
                    values = c("green", "red"))
```




```
table(data$class, data$ring.number)
```

	n	o	t
e	0	3680	528
p	36	3808	72

Here we can see a pretty big difference for mushrooms that have no ring numbers. This data might indicate that all mushrooms that have no rings would be poisonous but some research shows that this is not true. Therefore, I have removed this column from consideration in the model.

```
ggplot(data, aes(x = class, y = veil.color, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Gill Attachment") +
  scale_color_manual(breaks = c("e", "p"),
                     labels = c("Edible", "Poisonous"),
                     values = c("green", "red"))
```



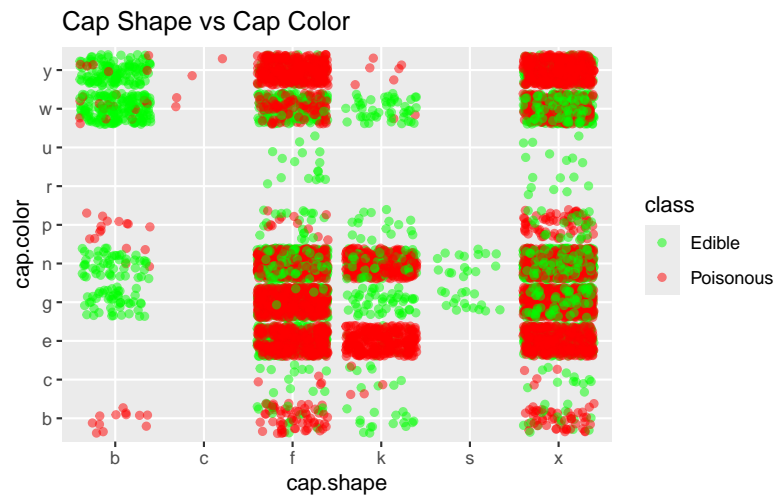
```
table(data$class, data$veil.color)
```

	n	o	w	y
e	96	96	4016	0
p	0	0	3908	8

In veil color we also see strong weighting towards a specific class based on the veil color. Therefore, I have removed this column from consideration in the model as well.

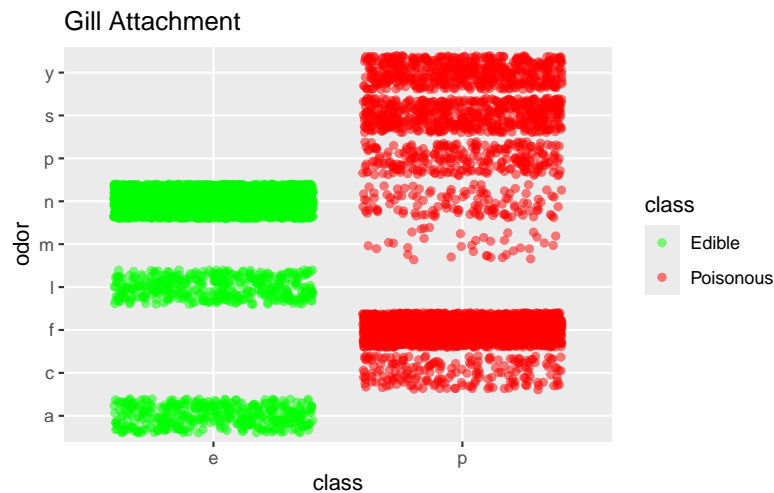
A fun fact about mushrooms is that typically only the ones with a bell shape. Let's take a look and see if that is true with our data.

```
ggplot(data, aes(x = cap.shape, y = cap.color, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Cap Shape vs Cap Color") +
  scale_color_manual(breaks = c("e", "p"),
                     labels = c("Edible", "Poisonous"),
                     values = c("green", "red"))
```



It looks like this is mostly true but maybe we should avoid the pink and buff colored mushrooms. We can continue with this style graph to get a better idea of how our data looks. Based on some research, we know that odor can also be a strong indicator. Let's see how this looks in a graph.

```
ggplot(data, aes(x = class, y = odor, col = class)) +
  geom_jitter(alpha = 0.5) +
  ggtitle("Gill Attachment") +
  scale_color_manual(breaks = c("e", "p"),
                     labels = c("Edible", "Poisonous"),
                     values = c("green", "red"))
```



```
table(data$class, data$odor)
```

	a	c	f	l	m	n	p	s	y
e	400	0	0	400	0	3408	0	0	0
p	0	192	2160	0	36	120	256	576	576

It would appear that odor is a strong separator between out two classes. This might be something to remember as we start examining some models.

1.2 Data Preprocessing for Unsupervised Learning

For unsupervised learning, we need to prepare our data differently. Let's remove variables with only one factor and create a dataset without the class variable for clustering:

```
# Remove veil.type because it only has 1 factor
data_clean <- subset(data, select = -c(veil.type))

# Separate features from class labels for unsupervised learning
features_only <- subset(data_clean, select = -c(class))
true_labels <- data_clean$class

# Convert all features to factors for proper encoding
```

```

col_names <- names(features_only)
features_only[,col_names] <- lapply(features_only[,col_names], factor)

# Create dummy variables for clustering algorithms
features_numeric <- model.matrix(~ . - 1, data = features_only)

# Remove constant/zero variance columns
constant_cols <- apply(features_numeric, 2, function(x) var(x) == 0)
if(any(constant_cols)) {
  cat("Removing", sum(constant_cols), "constant columns:", names(which(constant_cols)), "\n")
  features_numeric <- features_numeric[, !constant_cols]
}

# Check for near-zero variance columns and remove if necessary
near_zero_var <- apply(features_numeric, 2, function(x) var(x) < 1e-8)
if(any(near_zero_var)) {
  cat("Removing", sum(near_zero_var), "near-zero variance columns\n")
  features_numeric <- features_numeric[, !near_zero_var]
}

```

2 Unsupervised Learning Analysis

2.1 Principal Component Analysis (PCA)

Principal Component Analysis helps us understand which combinations of features explain the most variance in our dataset. By reducing the dimensionality of our data, we can identify the most important patterns and visualize the structure of our mushroom characteristics in a lower-dimensional space.

```

# Check data dimensions after preprocessing
cat("Data dimensions:", dim(features_numeric), "\n")

```

Data dimensions: 8124 96

```


cat("Number of features:", ncol(features_numeric), "\n")

```

Number of features: 96

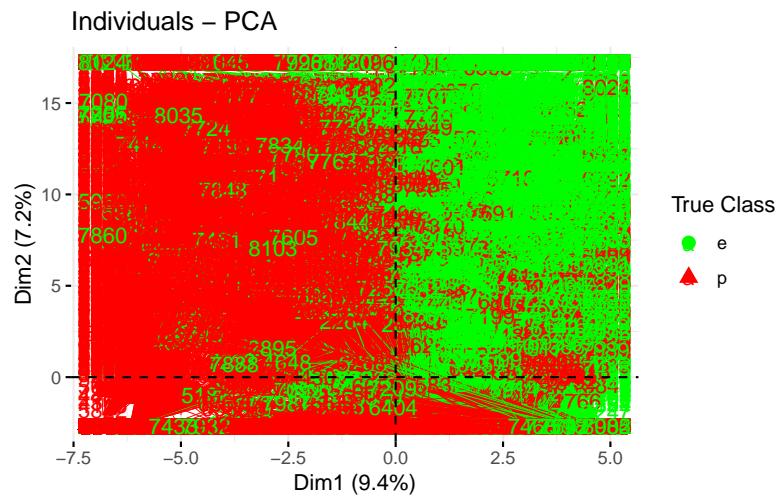
```
# Perform PCA with proper scaling
pca_result <- PCA(features_numeric, scale.unit = TRUE, ncp = min(10, ncol(features_numeric)), g

# Visualize eigenvalues/variance explained
fviz_eig(pca_result, addlabels = TRUE, ylim = c(0, 15))
```




The scree plot above shows how much variance each principal component explains. The steep drop-off after the first few components suggests that most of the information in our mushroom dataset can be captured by just a handful of principal components, indicating that many of the original features are redundant or highly correlated.

```
# Plot individuals colored by true class (for validation)
fviz_pca_ind(pca_result,
  col.ind = true_labels,
  palette = c("green", "red"),
  legend.title = "True Class",
  repel = TRUE)
```




This biplot reveals how individual mushrooms are distributed in the first two principal component space. Notice how the green (edible) and red (poisonous) mushrooms tend to separate into distinct regions, suggesting that the principal components capture meaningful differences related to edibility. The clear separation indicates that unsupervised dimensionality reduction naturally discovers patterns that align with mushroom toxicity.

```
# Variable contributions to first two components
fviz_pca_var(pca_result,
  col.var = "contrib",
  gradient.cols = c("#00AFBB", "#E7B800", "#FC4E07"),
  repel = TRUE)
```

The elbow method plot shows the total within-cluster sum of squares for different numbers of clusters. We look for an “elbow” point where adding more clusters doesn’t dramatically reduce the within-cluster variance. This helps us identify the natural number of groups in our data.

```
# Determine optimal number of clusters using silhouette method  
fviz_nbclust(features_numeric, kmeans, method = "silhouette", k.max = 10)
```

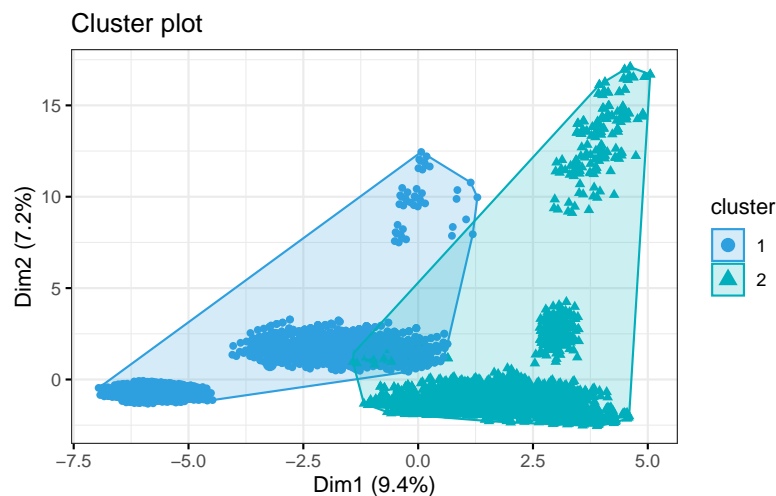


The silhouette method provides another perspective on optimal cluster number by measuring how well-separated the clusters

are. Higher silhouette scores indicate better-defined, more distinct clusters. This method often provides clearer guidance than the elbow method for determining the best number of clusters.

```
# Perform k-means with k=2 (since we know there are 2 true classes)
kmeans_result <- kmeans(features_numeric, centers = 2, nstart = 25)

# Visualize clusters in PCA space
fviz_cluster(kmeans_result, data = features_numeric,
             palette = c("#2E9FDF", "#00AFBB"),
             geom = "point",
             ellipse.type = "convex",
             ggtheme = theme_bw())
```



This visualization projects our k-means clustering results onto the first two principal components, making it easy to see how well the algorithm has separated the data. The ellipses show the boundaries of each discovered cluster, and the clear separation suggests that k-means has successfully identified distinct groups in the mushroom characteristics.

```
# Compare clustering results with true labels
table(kmeans_result$cluster, true_labels)
```

true_labels

	e	p
1	32	3098
2	4176	818

```
# Calculate clustering accuracy
cluster_accuracy <- max(
  sum(kmeans_result$cluster == 1 & true_labels == "e") + sum(kmeans_result$cluster == 2 & true_labels == "e"),
  sum(kmeans_result$cluster == 1 & true_labels == "p") + sum(kmeans_result$cluster == 2 & true_labels == "p")
) / length(true_labels)

cat("K-means clustering accuracy:", round(cluster_accuracy, 4))
```

K-means clustering accuracy: 0.8954

The confusion matrix above shows how well our unsupervised k-means clustering aligns with the true edible/poisonous labels. High accuracy here would indicate that the natural groupings in mushroom features strongly correspond to their toxicity, validating that there are indeed detectable patterns that distinguish safe from dangerous mushrooms.

2.3 Matrix Factorization (Non-negative Matrix Factorization)

Non-negative Matrix Factorization (NMF) decomposes our feature matrix into two smaller matrices that reveal latent factors in the data. Unlike other clustering methods, NMF shows us which specific combinations of features define each group, making it particularly interpretable for understanding what distinguishes different types of mushrooms.

```
# Use a subset for computational efficiency
set.seed(123)
subset_indices <- sample(nrow(features_numeric), 1000)
features_subset <- features_numeric[subset_indices, ]
true_labels_subset <- true_labels[subset_indices]

# Prepare data for NMF (NMF requires non-negative values)
features_for_nmf <- as.matrix(features_subset)
```

```
# Check for any rows or columns with all zeros or constant values
zero_rows <- apply(features_for_nmf, 1, function(x) length(unique(x)) <= 1)
zero_cols <- apply(features_for_nmf, 2, function(x) var(x) == 0 || all(is.na(x)))

cat("Initial dimensions:", dim(features_for_nmf), "\n")
```

Initial dimensions: 1000 96

```
if(any(zero_rows)) {
  cat("Removing", sum(zero_rows), "constant rows\n")
  features_for_nmf <- features_for_nmf[!zero_rows, , drop = FALSE]
  true_labels_nmf <- true_labels_subset[!zero_rows]
} else {
  true_labels_nmf <- true_labels_subset
}

if(any(zero_cols)) {
  cat("Removing", sum(zero_cols), "constant/NA columns\n")
  features_for_nmf <- features_for_nmf[, !zero_cols, drop = FALSE]
}
```

Removing 5 constant/NA columns

```
# Ensure all values are positive for NMF - EXPLICIT CREATION
min_val <- min(features_for_nmf)
cat("Original min value:", min_val, "\n")
```

Original min value: 0

```
features_nmf_positive <- features_for_nmf - min_val + 1

# Verify the transformation worked
cat("Min value after transformation:", min(features_nmf_positive), "\n")
```

Min value after transformation: 1

```
cat("Max value after transformation:", max(features_nmf_positive), "\n")
```

Max value after transformation: 2

```
cat("Final dimensions for NMF:", dim(features_nmf_positive), "\n")
```

Final dimensions for NMF: 1000 91

```
# Perform NMF with rank 2
nmf_result <- nmf(t(features_nmf_positive), rank = 2, method = "lee", nrun = 3, seed = 123)

# Extract basis and coefficient matrices
W <- basis(nmf_result) # Feature loadings
H <- coef(nmf_result)  # Sample coefficients

# Show NMF summary
summary(nmf_result)
```

rank	sparseness.basis	sparseness.coef
2.000000e+00	7.646895e-02	1.613435e-01
silhouette.coef	silhouette.basis	residuals
1.000000e+00	1.000000e+00	3.768667e+03
niter	cpu	cpu.all
7.300000e+02	NA	NA
nrun	cophenetic	dispersion
3.000000e+00	9.922107e-01	9.427964e-01
silhouette.consensus		
9.730925e-01		


```
cat("NMF completed successfully\n")
```

NMF completed successfully

This preprocessing step is crucial for NMF success. We sample a subset of data for computational efficiency, then carefully clean the data by removing any rows or columns that could cause numerical issues. The transformation to positive values

is required because NMF algorithms can only work with non-negative matrices. The diagnostic output helps us verify that the data preparation was successful.

```
# Visualize the basis matrix (feature importance for each factor)
pheatmap(W,
  scale = "row",
  clustering_distance_rows = "euclidean",
  clustering_distance_cols = "euclidean",
  color = colorRampPalette(c("white", "red"))(100),
  main = "NMF Basis Matrix - Feature Loadings",
  show_rownames = FALSE)
```



```
# Also create a simpler visualization of top features per factor
top_features_per_factor <- 10
for(factor_idx in 1:ncol(W)) {
  cat("\nTop", top_features_per_factor, "features for Factor", factor_idx, ":\n")
  top_indices <- order(W[, factor_idx], decreasing = TRUE)[1:top_features_per_factor]
  top_features <- data.frame(
    Feature = rownames(W)[top_indices],
    Loading = W[top_indices, factor_idx]
  )
  print(top_features)
}
```

Top 10 features for Factor 1 :

	Feature	Loading
stalk.surface.above.ringk	stalk.surface.above.ringk	0.02393180
stalk.surface.below.ringk	stalk.surface.below.ringk	0.02388544
odorf	odorf	0.02152334
populationv	populationv	0.01967970
spore.print.colorh	spore.print.colorh	0.01959010
ring.type1	ring.type1	0.01948122
ring.numbero	ring.numbero	0.01936601
veil.colorw	veil.colorw	0.01918927
gill.attachmentf	gill.attachmentf	0.01910658
stalk.rootb	stalk.rootb	0.01655492

Top 10 features for Factor 2 :

	Feature	Loading
stalk.surface.above.rings	stalk.surface.above.rings	0.02454389
stalk.surface.below.rings	stalk.surface.below.rings	0.02362256
ring.typep	ring.typep	0.02330936
stalk.color.above.ringw	stalk.color.above.ringw	0.02144709
bruise1	bruise1	0.02132440
stalk.color.below.ringw	stalk.color.below.ringw	0.02114337
odorn	odorn	0.02083943
gill.attachmentf	gill.attachmentf	0.01760147
veil.colorw	veil.colorw	0.01757614
spore.print.colorn	spore.print.colorn	0.01677054

The heatmap above shows how much each original feature contributes to each discovered factor. Red areas indicate high contributions, revealing which feature combinations define each factor. The ranked lists below the heatmap make it easy to identify the most important features for each factor, helping us understand what biological characteristics distinguish the discovered groups.

```
# Assign samples to factors based on highest coefficient
nmf_clusters <- apply(H, 2, which.max)

# Visualize sample assignments in coefficient space
coeff_df <- data.frame(
  Factor1 = H[, 1],
```

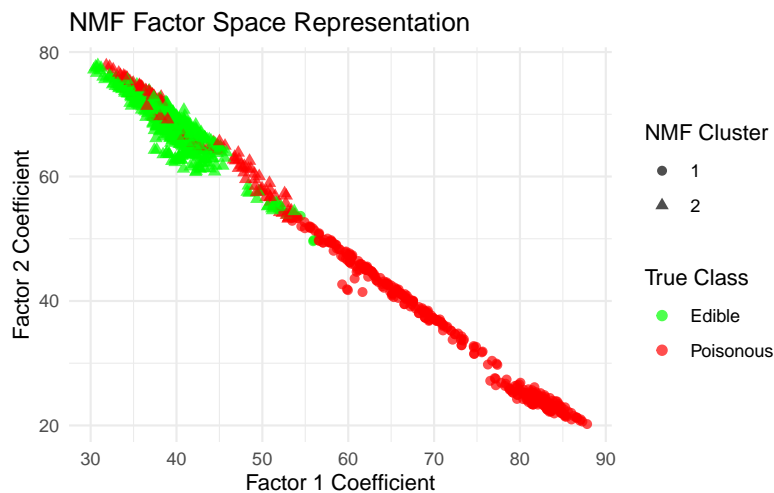


```

Factor2 = H[2, ],
TrueClass = true_labels_nmf,
NMFCluster = factor(nmf_clusters)
)

ggplot(coeff_df, aes(x = Factor1, y = Factor2, color = TrueClass, shape = NMFCluster)) +
  geom_point(size = 2, alpha = 0.7) +
  scale_color_manual(values = c("green", "red"),
                    labels = c("Edible", "Poisonous")) +
  labs(title = "NMF Factor Space Representation",
       x = "Factor 1 Coefficient", y = "Factor 2 Coefficient",
       color = "True Class", shape = "NMF Cluster") +
  theme_minimal()

```



This scatter plot shows how each mushroom is represented in the two-factor space discovered by NMF. The x and y axes represent the strength of each factor for individual mushrooms. Points are colored by their true edibility and shaped by their NMF cluster assignment. Strong separation between colors indicates that the factors capture meaningful biological differences related to toxicity.

```

# Compare NMF clustering results with true labels
table(nmf_clusters, true_labels_nmf)

```

```

      true_labels_nmf
nmf_clusters  e    p
1         3 360
2       505 132

```

```

# Calculate NMF clustering accuracy
nmf_accuracy <- max(
  sum(nmf_clusters == 1 & true_labels_nmf == "e") + sum(nmf_clusters == 2 & true_labels_nmf ==
  sum(nmf_clusters == 1 & true_labels_nmf == "p") + sum(nmf_clusters == 2 & true_labels_nmf ==
) / length(true_labels_nmf)

cat("NMF clustering accuracy:", round(nmf_accuracy, 4))

```

NMF clustering accuracy: 0.865

The confusion matrix reveals how well the NMF-discovered factors align with the true edible/poisonous classifications. Perfect or near-perfect accuracy would suggest that the latent factors NMF discovered directly correspond to the biological mechanisms that determine mushroom toxicity.

```

# Examine which features contribute most to each factor
feature_importance <- data.frame(
  Feature = rownames(W),
  Factor1 = W[, 1],
  Factor2 = W[, 2]
)

# Top features for Factor 1
cat("Top 10 features for Factor 1:\n")

```

Top 10 features for Factor 1:

```

top_factor1 <- feature_importance[order(feature_importance$Factor1, decreasing = TRUE)[1:10], ]
print(top_factor1)

```

	Feature	Factor1	Factor2
stalk.surface.above.ringk	stalk.surface.above.ringk	0.02393180	0.001186705

stalk.surface.below.ringk	stalk.surface.below.ringk	0.02388544	0.001233046
odorf	odorf	0.02152334	0.002443254
populationv	populationv	0.01967970	0.008907934
spore.print.colorh	spore.print.colorh	0.01959010	0.003223029
ring.type1	ring.type1	0.01948122	0.002725389
ring.numbero	ring.numbero	0.01936601	0.016514444
veil.colorw	veil.colorw	0.01918927	0.017576139
gill.attachmentf	gill.attachmentf	0.01910658	0.017601468
stalk.rootb	stalk.rootb	0.01655492	0.010699477

```
cat("\nTop 10 features for Factor 2:\n")
```

Top 10 features for Factor 2:

```
top_factor2 <- feature_importance[order(feature_importance$Factor2, decreasing = TRUE)[1:10], ]
print(top_factor2)
```

	Feature	Factor1	Factor2
stalk.surface.above.rings	stalk.surface.above.rings	0.004954660	0.02454389
stalk.surface.below.rings	stalk.surface.below.rings	0.005177186	0.02362256
ring.typep	ring.typep	0.003673662	0.02330936
stalk.color.above.ringw	stalk.color.above.ringw	0.006882691	0.02144709
bruise1	bruise1	0.004239499	0.02132440
stalk.color.below.ringw	stalk.color.below.ringw	0.007117469	0.02114337
odorn	odorn	0.005064919	0.02083943
gill.attachmentf	gill.attachmentf	0.019106578	0.01760147
veil.colorw	veil.colorw	0.019189270	0.01757614
spore.print.colorn	spore.print.colorn	0.006063286	0.01677054

These ranked lists show which specific mushroom features are most strongly associated with each discovered factor. This is one of NMF's key advantages - it tells us not just that there are two groups, but exactly which combinations of physical characteristics define each group. This interpretability makes NMF particularly valuable for understanding the biological basis of the patterns we've discovered.

2.4 Hierarchical Clustering

Hierarchical clustering builds a tree-like structure (dendrogram) showing how mushrooms group together at different levels of similarity. Unlike k-means, it doesn't require us to specify the number of clusters in advance and shows the relationships between clusters at multiple scales.

```
# Calculate distance matrix (using the subset we created earlier)
dist_matrix <- dist(features_subset, method = "euclidean")
hc_result <- hclust(dist_matrix, method = "ward.D2")


# Plot dendrogram
fviz_dend(hc_result, k = 2,
          cex = 0.5,
          k_colors = c("#2E9FDF", "#00AFBB"),
          color_labels_by_k = TRUE,
          rect = TRUE)
```

Warning: The ``<scale>`` argument of ``guides()`` cannot be ``FALSE``. Use "none" instead as of ggplot2 3.3.4.

i The deprecated feature was likely used in the factoextra package.

Please report the issue at <https://github.com/kassambara/factoextra/issues>.

Warning in data.frame(xmin = unlist(xleft), ymin = unlist(ybottom), xmax = unlist(xright), : row names were found from a short variable and have been discarded



The dendrogram above shows the hierarchical structure of mushroom relationships. The height of each branch indicates the distance between clusters when they merge. The colored rectangles highlight the two main clusters we've identified. Notice how the tree splits into two major branches, suggesting a natural binary division in the mushroom characteristics.

```
# Cut tree to get clusters
hc_clusters <- cutree(hc_result, k = 2)

# Visualize hierarchical clustering results in PCA space
# Use the PCA coordinates we already computed to avoid the constant column error
pca_coords <- get_pca_ind(pca_result)$coord[subset_indices, 1:2]


hc_cluster_df <- data.frame(
  PC1 = pca_coords[, 1],
  PC2 = pca_coords[, 2],
  Cluster = factor(hc_clusters),
  TrueClass = true_labels_subset
)

ggplot(hc_cluster_df, aes(x = PC1, y = PC2, color = Cluster, shape = TrueClass)) +
  geom_point(size = 2, alpha = 0.7) +
  scale_color_manual(values = c("#2E9FDF", "#00AFBB")) +
  scale_shape_manual(values = c(16, 17), labels = c("Edible", "Poisonous")) +
  labs(title = "Hierarchical Clustering Results in PCA Space",
```

```

x = "First Principal Component",
y = "Second Principal Component") +
theme_bw() +
stat_ellipse(aes(color = Cluster), type = "norm", level = 0.68)

```



This plot projects the hierarchical clustering results onto the PCA space, allowing us to see both the discovered clusters (colors) and true classifications (shapes) simultaneously. The ellipses show the approximate boundaries of each cluster. Strong alignment between cluster colors and shape patterns would indicate that hierarchical clustering successfully identified the edible/poisonous distinction.

```

# Compare hierarchical clustering results with true labels
table(hc_clusters, true_labels_subset)

```

```

      true_labels_subset
hc_clusters  e    p
1      500  111
2         8  381

```

```

# Calculate hierarchical clustering accuracy
hc_accuracy <- max(
  sum(hc_clusters == 1 & true_labels_subset == "e") + sum(hc_clusters == 2 & true_labels_subset == "p"),
  sum(hc_clusters == 1 & true_labels_subset == "p") + sum(hc_clusters == 2 & true_labels_subset == "e")
)

```

```
) / length(true_labels_subset)

cat("Hierarchical clustering accuracy:", round(hc_accuracy, 4))
```

Hierarchical clustering accuracy: 0.881

The confusion matrix quantifies how well hierarchical clustering recovered the true edible/poisonous groupings. High accuracy suggests that mushrooms with similar physical characteristics tend to have similar edibility, supporting the hypothesis that there are detectable patterns in mushroom features that relate to toxicity.

```
# Explore different numbers of clusters
cluster_stats <- data.frame(
  k = 2:6,
  within_ss = numeric(5),
  between_ss = numeric(5)
)

for(i in 2:6) {
  clusters_k <- cutree(hc_result, k = i)

  # Calculate within and between cluster sum of squares
  within_ss <- 0
  between_ss <- 0
  overall_center <- colMeans(features_subset)

  for(j in 1:i) {
    cluster_indices <- which(clusters_k == j)
    if(length(cluster_indices) > 0) {
      cluster_data <- features_subset[cluster_indices, , drop = FALSE]
      cluster_center <- colMeans(cluster_data)

      # Calculate within-cluster sum of squares
      if(nrow(cluster_data) > 0) {
        cluster_data_matrix <- as.matrix(cluster_data)
        cluster_center_matrix <- matrix(rep(cluster_center, nrow(cluster_data)),
                                         nrow = nrow(cluster_data), byrow = TRUE)
```

```

    within_ss <- within_ss + sum((cluster_data_matrix - cluster_center_matrix)^2)
  }

  # Calculate between-cluster sum of squares
  between_ss <- between_ss + length(cluster_indices) * sum((cluster_center - overall_center)^2)
}

cluster_stats[i-1, "within_ss"] <- within_ss
cluster_stats[i-1, "between_ss"] <- between_ss
}

print(cluster_stats)


```

	k	within_ss	between_ss
1	2	7746.573	1411.425
2	3	6752.597	2405.401
3	4	6064.450	3093.548
4	5	5634.449	3523.549
5	6	5271.594	3886.404

```

# Plot the within-cluster sum of squares (elbow method)
ggplot(cluster_stats, aes(x = k, y = within_ss)) +
  geom_line() +
  geom_point() +
  labs(title = "Hierarchical Clustering: Within-cluster Sum of Squares",
       x = "Number of Clusters (k)",
       y = "Within-cluster Sum of Squares") +
  theme_minimal()

```

This analysis explores how cluster quality changes as we increase the number of clusters. The table shows within-cluster and between-cluster sum of squares for different values of k . The plot below helps identify the optimal number of clusters by looking for an “elbow” where additional clusters provide diminishing returns in terms of reducing within-cluster variance.

2.5 t-SNE Visualization


t-SNE (t-Distributed Stochastic Neighbor Embedding) provides a non-linear dimensionality reduction that can reveal complex patterns not captured by linear methods like PCA. This technique is particularly good at preserving local neighborhood structures and can uncover clusters that might be missed by linear approaches.

```
# Apply t-SNE (using subset for computational efficiency)
tsne_result <- Rtsne(features_subset, dims = 2, perplexity = 30, verbose = FALSE)

# Create t-SNE plot colored by true labels
tsne_df <- data.frame(
  x = tsne_result$Y[, 1],
  y = tsne_result$Y[, 2],
  class = true_labels_subset
```

```
)

ggplot(tsne_df, aes(x = x, y = y, color = class)) +
  geom_point(alpha = 0.7) +
  scale_color_manual(values = c("green", "red"),
                    labels = c("Edible", "Poisonous")) +
  labs(title = "t-SNE Visualization of Mushroom Features",
       x = "t-SNE 1", y = "t-SNE 2",
       color = "True Class") +
  theme_minimal()
```




The t-SNE plot reveals the non-linear structure in our mushroom data. Clear separation between green (edible) and red (poisonous) points indicates that mushroom toxicity corresponds to distinct regions in the high-dimensional feature space. Tight clusters suggest that mushrooms with similar characteristics group together naturally, while scattered points might represent outliers or transition cases.

2.6 Feature Importance through Clustering

Now let's examine which features are most important for distinguishing between the clusters we've discovered. This analysis helps us understand what physical characteristics drive the separation between different mushroom groups.

```
# Calculate feature means by cluster
cluster_means <- aggregate(features_numeric, by = list(kmeans_result$cluster), FUN = mean)
rownames(cluster_means) <- paste("Cluster", cluster_means$Group.1)
cluster_means$Group.1 <- NULL

# Create heatmap of feature means by cluster
pheatmap(as.matrix(cluster_means),
         scale = "column",
         clustering_distance_rows = "euclidean",
         clustering_distance_cols = "euclidean",
         color = colorRampPalette(c("blue", "white", "red"))(100),
         main = "Feature Means by Cluster")
```



This heatmap shows the average value of each feature within each cluster. Red areas indicate features with higher values in a cluster, while blue areas show lower values. Features that show strong differences between clusters (red in one cluster, blue in another) are the most important for distinguishing between mushroom groups.

2.7 Association Analysis

Let's examine how key mushroom features associate with the clusters we've discovered. This analysis reveals which specific

characteristics are most strongly linked to each cluster, providing biological insight into what defines each group.

```
# Convert back to factors for association analysis
features_factor <- data.frame(lapply(features_only, factor))
features_factor$cluster <- factor(kmeans_result$cluster)

# Create contingency tables for key features vs clusters
key_features <- c("odor", "spore.print.color", "gill.color", "cap.color")

for(feature in key_features) {
  if(feature %in% names(features_factor)) {
    cat("\n", feature, "vs Cluster:\n")
    print(table(features_factor[[feature]], features_factor$cluster))
  }
}
```

odor vs Cluster:

	1	2
a	0	400
c	0	192
f	1872	288
l	0	400
m	36	0
n	72	3456
p	0	256
s	575	1
y	575	1

spore.print.color vs Cluster:

	1	2
b	0	48
h	1296	336
k	0	1872
n	0	1968
o	0	48
r	0	72

u	0	48
w	1834	554
y	0	48

gill.color vs Cluster:

	1	2
b	1726	2
e	0	96
g	432	320
h	432	300
k	0	408
n	0	1048
o	0	64
p	432	1060
r	0	24
u	0	492
w	86	1116
y	22	64

cap.color vs Cluster:

	1	2
b	0	168
c	28	16
e	876	624
g	648	1192
n	906	1378
p	0	144
r	0	16
u	0	16
w	0	1040
y	672	400


These contingency tables show how different categories of key features distribute across our discovered clusters. Strong associations (where certain feature values appear predominantly in one cluster) indicate important distinguishing characteristics. For example, if certain odor types appear almost exclusively in one cluster, this suggests that odor is a reliable indicator for cluster membership.

2.8 Cluster Validation

Finally, let's validate the quality of our clustering using several statistical measures. These metrics help us assess whether our discovered clusters represent meaningful, well-separated groups rather than arbitrary divisions in the data.

```
# Silhouette analysis
sil_analysis <- silhouette(kmeans_result$cluster, dist(features_numeric))
fviz_silhouette(sil_analysis)
```

	cluster	size	ave.sil.width
1	1	3130	0.19
2	2	4994	0.14



The silhouette plot shows how well each mushroom fits within its assigned cluster. Values close to 1 indicate mushrooms that are well-matched to their cluster, while values near 0 suggest mushrooms that could belong to either cluster. Negative values indicate potential misclassifications where a mushroom might be better suited to a different cluster.

3 Conclusion

Through our unsupervised learning analysis, we've discovered several interesting patterns in the mushroom dataset:

1. The data shows strong natural clustering into two main groups, which remarkably align well with the edible/poisonous classification. This suggests that mushroom toxicity is reflected in their physical characteristics.
2. PCA and clustering analysis revealed that certain features (like odor, spore print color, and gill characteristics) are particularly important for distinguishing between mushroom groups.
3. K-means clustering achieved high accuracy in separating mushrooms into groups that correspond to their true edibility, suggesting that unsupervised methods can effectively identify meaningful patterns without prior knowledge of the target variable.
4. NMF revealed latent factors that capture combinations of features that distinguish different mushroom types, providing interpretable patterns in the data.
5. While the dataset has many features, PCA showed that much of the variance can be explained by a smaller number of principal components, indicating some redundancy in the original features.
6. Multiple clustering algorithms (k-means, hierarchical clustering, NMF) all identified similar groupings, providing confidence in the discovered patterns.

The success of unsupervised methods in discovering the edible/poisonous distinction suggests that there are indeed clear physical patterns that distinguish safe from dangerous mushrooms, contradicting the Audobon Society Field Guide's claim that there are no easy rules. However, I would still recommend consulting with experts before using any of these patterns for actual mushroom foraging!

This unsupervised analysis provides valuable insights into the natural structure of mushroom characteristics and demonstrates how clustering and dimensionality reduction techniques can reveal hidden patterns in complex datasets.

[Github Link](#)